

EDUCATION

George Mason University, Virginia

August 2016 – May 2018

Master of Science, Computer Science

Relevant coursework: Theory & Applications of Data Mining, Data Mining on Multimedia Data, Pattern Recognition, Applied Statistics, Mining Massive Datasets, Analysis of Algorithms, Distributed Systems

University of Mumbai, India

May 2010 – June 2014

Bachelor of Engineering, Computer Engineering

SKILLS

Languages: Python (3 years), Java (4+ years), SQL(4+ years), C++ (< 1 year), R (< 1 year)

Tools: Tableau, IPython, Weka, Docker, AWS, Git, Apache Airflow, CI/CD (Travis, CircleCI)

Techniques: Clustering, classification, natural language processing, regression, matrix factorization

Libraries: scikit-learn, Numpy, Pandas, Scipy, GraphLab, matplotlib, seaborn, gensim

Beginner level knowledge of Hadoop, Spark, Tensorflow and Keras.

PROJECTS

Music Mood Classification Using the Million Song Dataset: Used audio features and machine learning for mood analysis. Techniques: segment aggregation, XGBoost, Random Forest, Support Vector Machines

Quora Question Pairs: A Kaggle competition about detecting duplicate questions on Quora. Achieved a log loss of 0.324 on the leaderboard using an LSTM built with Keras and TensorFlow.

Movie Recommender System: Predicted movie ratings given by a user through item based collaborative filtering. Achieved an RMSE of 0.76. Techniques: factorization machines

Scipy: I've made a few contributions to the open source library and the closed PRs are visible here:

bit.do/scipy-prs

Drug Activity Prediction: Achieved an F1 score of 0.73 while classifying drugs as active/inactive in an imbalanced dataset. Techniques: Stochastic Gradient Descent (SGD)

Raven: An open-source CLI tool to manage your Spotify music library.

WORK EXPERIENCE

Software Engineer Intern- Data, Udacity Mountain View, CA

June 2017 – August 2017

- Improved organizational access to data by building a Slack bot with Flask and SQLAlchemy.
- Built an ETL system to migrate platform events to Amplitude for real-time analytics.
- Communicated insights about Udacity products after analyzing 300,000+ events.
- Built data pipelines using Apache Airflow.

GIS Programmer, George Mason University

August 2016 – present

- Developed George Mason University's OpenGeoportal using Java, Apache Solr, Python.
- Lead workshops on ArcGIS, QGIS, Python and CartoDB to train students and faculty.
- Assisted students with framing research questions and designing projects based around GIS and data analysis.

Senior Software Engineer, Capgemini

June 2014 – July 2016

- Developed a Java/J2EE application for a Fortune 500 engineering client.
- Built a prototype for a product lifecycle management application using Java.
- Improved the performance of SQL stored procedures by 0.5-2x.
- Developed features for forecasting, asset management and resource allocation in a web application using C# and ASP .NET.
- Improved perceived page load performance on dashboards by 30%.
- Reduced page load speed by 6s for frequently accessed pages using caching mechanisms.
- Mentored and onboarded 2 junior engineers by training them on the technical and business aspects of the project.